# DJ Link Packet Analysis

James Elliott
Deep Symmetry, LLC

May 28, 2017

### Abstract

The protocol used by Pioneer professional DJ equipment to communicate and coordinate performances can be monitored to provide useful information for synchronizing other software, such as light shows and sequencers. By creating a "virtual CDJ" that sends appropriate packets to the network, other devices can be induced to send packets containing even more useful information about their state. This article documents what has been learned so far about the protocol, and how to accomplish these tasks.

# Contents

```
      0   1   2   3   4   5   6   7   8   9   a   b   c   d   e   f
```

Header $\Big\{$

| 51 73 70 74 31 57 6d 4a 4f 4c **0a** 00 |
|---|
| Device Name (padded with 00) |

$\Big\}$ Name

| 01 02 00 25 02 |
|---|

Figure 1: Initial announcement packets from Mixer

```
      0   1   2   3   4   5   6   7   8   9   a   b   c   d   e   f
```

Header $\Big\{$

| 51 73 70 74 31 57 6d 4a 4f 4c **00** 00 |
|---|
| Device Name (padded with 00) |

$\Big\}$ Name

| 01 02 00 2c $N$ 02 | MAC address |
|---|---|

Figure 2: First-stage Mixer device number assignment packets

# 1   Mixer Startup

When the mixer starts up, after it obtains an IP address (or gives up on doing that and self-assigns an address), it sends out what look like a series of packets[1] simply announcing its existence to UDP port 50000 on the broadcast address of the local network.

These have a data length[2] of 37 bytes, appear roughly every 300 milliseconds, and have the content shown in Figure 1.

The tenth byte (inside what is labeled the header) is bolded because its value changes in the different types of packets which follow.

After about three of these packets are sent, another series of three begins. It is not clear what purpose these packets serve, because they are not yet asserting ownership of any device number; perhaps they are used when CDJs are powering up as part of the mechanism the mixer can use to tell them which device number to use based on which network port they are connected to?

In any case, these three packets have a data length of 44 bytes, are again sent to UDP port 50000 on the local network broadcast address, at roughly 300 millisecond intervals, and have the content shown in Figure 2.

The value $N$ at byte 36 is 1, 2, or 3, depending on whether this is

---

[1]The packet capture described in this analysis can be found at `https://github.com/brunchboy/dysentery/raw/master/doc/assets/powerup.pcapng`

[2]Values within packets are shown in hexadecimal, while packet lengths and byte offsets are discussed in decimal.

the first, second, or third time the packet is sent.

After these comes another series of three numbered packets. These appear to be claiming the device number for a particular device, as well as announcing the IP address at which it can be found. They have a data length of 50 bytes, and are again sent to UDP port 50000 on the local network broadcast address, at roughly 300 millisecond intervals, with the content shown in Figure 3.



Figure 3: Second-stage Mixer device number assignment packets

I identify these as claiming/identifying the device number because the value $D$ at byte 46 is the same as the device number that the mixer uses to identify itself (`0x21`) and the same is true for the corresponding packets seen from the CDJs (they use device numbers 2 and 3, as they are connected to those ports/channels on the mixer).

As with the previous series of three packets, the value $N$ at byte 47 takes on the values 1, 2, and 3 in the three packets.

These are followed by another three packets, perhaps the last stage of claiming the device number, again at 300 millisecond intervals, to the same port 50000. These shorter packets have 38 bytes of data and the content shown in Figure 4.

As before the value $D$ at byte 36 is the same as the device number that the mixer uses to identify itself (`0x21`) and $N$ at byte 37 takes on the values 1, 2, and 3 in the three packets.

Once those are sent, the mixer seems to settle down and send what looks like a keep-alive packet to retain presence on the network and



Figure 4: Final-stage Mixer device number assignment packets

ownership of its device number, at a less frequent interval. These packets are 54 bytes long, again sent to port 50000 on the local network broadcast address, roughly every second and a half. They have the content shown in Figure 5.

|   | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

Header { 51 73 70 74 31 57 6d 4a 4f 4c **06** 00 } Name
Device Name (padded with 00)
01 02 00 36 *D* 02   MAC address   IP address
01 00 00 00 02 00

Figure 5: Mixer keep-alive packets

## 2   CDJ Startup

When a CDJ starts up the procedure and packets are nearly identical, with groups of three packets sent at 300 millisecond intervals to port 50000 of the local network broadcast address. The only difference between Figure 6 and Figure 1 is the final byte, which is 0x01 for the CDJ, and was 0x02 for the mixer.

|   | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

Header { 51 73 70 74 31 57 6d 4a 4f 4c **0a** 00 } Name
Device Name (padded with 00)
01 02 00 25 01

Figure 6: Initial announcement packets from CDJ

Similarly, the next series of three packets from the CDJ are nearly identical to those from the mixer. The only difference between Figure 7 and Figure 2 is byte 37 (immediately after the packet counter $N$), which again is 0x01 for the CDJ, and was 0x02 for the mixer.

However it appears that in this capture the CDJ skips the second stage of claiming a device number, probably because it is configured to be automatically assigned a device number based on the port of the mixer to which it is connected, and we cannot see a packet that the mixer sent it assigning it that device number. Instead, it jumps right to the end of the third and final stage, sending a single 38-byte packet

|   | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

Header { 51 73 70 74 31 57 6d 4a 4f 4c **00** 00

Device Name (padded with 00) } Name

01 02 00 2c $N$ 01 | MAC address

Figure 7: First-stage CDJ device number assignment packets

with header byte 10 set to tt 04 (which identified the three packets of the third stage when the mixer was starting up), with content identical to Figure 4.

Even though the value of $N$ is 01, this is the only packet in this series that the CDJ sends. It would probably behave differently if configured to assign its own device number (behaving like we saw the mixer behave in claiming its device number).

The CDJ then moves to the keep-alive stage, sending out 54-byte packets with the content shown in Figure 8.

|   | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

Header { 51 73 70 74 31 57 6d 4a 4f 4c **06** 00

Device Name (padded with 00) } Name

01 02 00 36 $D$ 01 | MAC address | IP address

01 00 00 00 01 00

Figure 8: CDJ keep-alive packets

As seems to always be the case when comparing mixer and CDJ packets, the difference between this and Figure 5 is that byte 37 (following the device number $D$) has the value 01 rather than 02, and the same is true of the second-to-last byte in each of the packets. (Byte 52 is 01 in Figure 8 and 02 in Figure 5.

## 3   Tracking BPM and Beats

For some time now, Afterglow[3] has been able to synchronize its light shows with music being played on Pioneer equipment by observing packets broadcast by the mixer to port 50001. Until recently, however, it was not possible to tell which player was the master, so there was

---

[3]https://github.com/brunchboy/afterglow#afterglow

|    | 0  | 1  | 2  | 3  | 4  | 5  | 6  | 7  | 8  | 9  | a  | b  | c  | d  | e  | f  |
|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 00 | 51 | 73 | 70 | 74 | 31 | 57 | 6d | 4a | 4f | 4c | **28** |    |    |    |    |    |
| 10 | \multicolumn Device Name (padded with 00) | | | | | | | | | | | | | | | 01 |
| 20 | 00 | $D$ | 00 | 3c | \multicolumn $nextBeat$ | | | | $2ndBeat$ | | | | $nextBar$ | | | |
| 30 | \multicolumn $4thBeat$ | | | | $2ndBar$ | | | | $8thBeat$ | | | | ff | ff | ff | ff |
| 40 | ff | ff | ff | ff | ff | ff | ff | ff | ff | ff | ff | ff | ff | ff | ff | ff |
| 50 | ff | ff | ff | ff | \multicolumn $Pitch$ | | | | 00 | 00 | $BPM$ | | $B_b$ | 00 | 00 | $D$ |

Figure 9: Beat packets

no way to determine the down beat (the start of each measure). Now that it is possible to to determine which CDJ is the master player using the packets described in Section 4, these beat packets have become far more useful, and Afterglow will soon be using them to track the down beat based on the beat number reported by the master player.

To track beats, open a socket and bind it to port 50001. The devices seem to broadcast two different kinds of packets to this port, a shorter packet containing 45 bytes of data, and a longer packet containing 96 bytes. The shorter packets seem to all have identical content, and do not seem to convey useful information, so we currently simply ignore them.

The 96-byte packets are sent on each beat, so even the arrival of the packet is interesting information, it means that the player is starting a new beat. (CDJs send these packets only when they are playing *and only for rekordbox-analyzed tracks.* The mixer sends them all the time, acting as a backup metronome when no other device is counting beats.) The content of these packets is shown in Figure 9.

The Device Number in $D$ (bytes 33 and 95) is the Player Number as displayed on the CDJ itself, or 33 for the mixer, or another value for a computer running rekordbox.

To facilitate synchronization of variable-tempo tracks, the number of milliseconds after which a variety of upcoming beats will occur are reported. *nextBeat* at bytes 36–39 is the number of milliseconds in which the very next beat will arrive. *2ndBeat* (bytes 40–43) is the number of milliseconds until the beat after that. *nextBar* (bytes 44–47) reports the number of milliseconds until the next measure of music begins, which may be from 1 to 4 beats away. *4thBeat* (bytes 48–51) reports how many millisecond will elapse until the fourth upcoming beat; *2ndBar* (bytes 52–55) the interval until the second measure after the current one begins (which will occur in 5 to 8 beats, depending how far

into the current measure we have reached); and $8thBeat$ (bytes 56–59) tells how many millieconds we have to wait until the eighth upcoming beat will arrive.

The player's current pitch adjustment[4] can be found in bytes 84–87, labeled $Pitch$. It represents a three-byte pitch adjustment percentage, where `0x00100000` represents no adjustment (0%), `0x00000000` represents slowing all the way to a complete stop ($-100\%$, reachable only in Wide tempo mode), and `0x00200000` represents playing at double speed ($+100\%$).

The pitch adjustment percentage represented by $Pitch$ is calculated as follows:

$$100 \times \frac{(byte[85] \times 65536 + byte[86] \times 256 + byte[87]) - 1048576}{1048576}$$

The current BPM of the track playing on the device[5] can be found at bytes 90–91 (labeled $BPM$). It is a two-byte integer representing one hundred times the current track BPM. So, the current track BPM value to two decimal places can be calculated as:

$$\frac{byte[90] \times 256 + byte[91]}{100}$$

In order to obtain the actual playing BPM (the value shown in the BPM display), this value must be multiplied by the current pitch adjustment. Since calculating the effective BPM reported by a CDJ is a common operation, here a simplified equation that results in the effective BPM to two decimal places, by combinining the $BPM$ and $Pitch$ values:[6]

$$\frac{(byte[90] \times 256 + byte[91]) \times (byte[85] \times 65536 + byte[86] \times 256 + byte[87])}{104857600}$$

The counter $B_b$ at byte 92 counts out the beat within each bar, cycling $1 \rightarrow 2 \rightarrow 3 \rightarrow 4$ repeatedly, and can be used to identify the down beat if it is coming from the master player.

## 4   Creating a Virtual CDJ

Although some useful information can be obtained simply by watching broadcast traffic on a network containing Pioneer gear, in order to

---

[4]The mixer always reports a pitch of $+0\%$.

[5]The mixer passes along the BPM of the master player.

[6]Since the mixer always reports a pitch adjustment of $+0\%$, its $BPM$ value can be used directly without this additional step.

get important details it is necessary to cause the gear to send you information directly. This can be done by simulating a "Virtual CDJ".[7]

To do this, bind a UDP server socket to port 50002 on the network interface on which you are receiving DJ-Link traffic, and start sending keep-alive packets to port 50000 on the broadcast address as if you were a CDJ. Follow the structure shown in Figure 8, but use the actual MAC and IP addresses of the network interface on which you are receiving DJ-Link traffic, so the devices can see how to reach you.

You can use a value like 5 for $D$ (the device/player number) so as not to conflict with any actual players you have on the network, and any name you would like. As long as you are sending these packets roughly every 1.5 seconds, the other players and mixers will begin sending packets directly to the socket you have opened on port 50002.

Each device seems to send status packets roughly every 200 milliseconds.

We are just beginning to analyze all the information which can be gleaned from these packets, but here is what we know so far.[8]

## 4.1   Mixer Status Packets

Packets from the mixer will have a length of 56 bytes and the content shown in Figure 10.

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| **00** | 51 | 73 | 70 | 74 | 31 | 57 | 6d | 4a | 4f | 4c | **29** | | | | | |
| **10** | Device Name (padded with 00) | | | | | | | | | | | | | | | 01 |
| **20** | 00 | $D$ | 00 | 14 | $D$ | 00 | 00 | $F$ | 00 | 10 | 00 | 00 | 80 | 00 | $BPM$ | |
| **30** | 00 | 10 | 00 | 00 | 00 | 09 | $X$ | $B_b$ | | | | | | | | |

Figure 10: Mixer status packets

Packets coming from a DJM-2000 nexus connected as the only mixer on the network contain a value of 33 (0x21) for their Device Number $D$ (bytes 33 and 36).

The value marked $F$ at byte 39 is evidently a status flag equivalent to the one shown in Figure 12, although on a mixer the only two values seen so far are 0xf0 when it is the tempo master, and 0xd0 when it

---

[7]Thanks are due to Diogo Santos for discovering the trick of creating a virtual CDJ in order to receive detailed status information from other devices.

[8]Examples of packets discussed in this section can be found in the capture at https://github.com/brunchboy/dysentery/raw/master/doc/assets/to-virtual.pcapng

is not. So evidently the mixer always considers itself to be playing and synced, but never on-air.

There are two places that might contain pitch values, bytes 41–43 and bytes 50–51, but since they always `0x100000` (or +0%), we can't be sure. The current tempo in beats-per-minute identified by the mixer can be obtained as:

$$\frac{byte[46] \times 256 + byte[47]}{100}$$

This value is labeled $BPM$ in Figure 10. Unfortunately, this BPM seems to only be valid when a rekordbox-analyzed source is playing; when the mixer is doing its own beat detection from unanalyzed audio sources, even though it displays the detected BPM on the mixer itself, and uses that to drive its beat effects, it does not send that value in these packets.

The current beat number within a bar (1, 2, 3 or 4) is sent in $byte[55]$, labeled $B_b$. However, the beat number is *not* synchronized with the master player, and these packets do not arrive at the same time as the beat started anyway, so this value is not useful for much. The beat number should be determined, when needed, from beat packets (described in Section 3) that are sent by the master player.

The value at $byte[54]$, labeled $X$, has an unknown meaning. It seems to start out with the value `0x00`, and then change when a player starts playing to the value `0xff`, but it may well do other things as well.

## 4.2 CDJ Status Packets

Packets from a CDJ will have a length of 212 bytes and the content shown in Figure 11 for nexus players. Older players send 208-byte packets with slightly less information. Newer firmware and Nexus 2 players send packets that are 284 or 292 bytes long.

The Device Number in $D$ (bytes 33 and 36) is the Player Number as displayed on the CDJ itself. In the case of this capture, the CDJs were assigned Player Numbers 2 and 3.

The activity flag $A$ at byte 39 seems to be 0 when the player is idle, and 1 when it is playing, searching, or loading a track.

When a rekordbox track is loaded, the device holding the rekordbox database from which the track was loaded is reported in $D_r$ at byte 40 (if the track was loaded from the local device, this will be the same as $D$; if it was loaded over the Link, it will be the number of a different device) When no track is loaded, $D_r$ has the value `00`.

Similarly, $S_r$ at byte 41 reports the slot from which the track was loaded: The value `00` means no track is loaded, `01` means the CD drive, `02` means the SD slot, and `03` means the USB slot. When a

|    | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|----|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| **00** | 51 | 73 | 70 | 74 | 31 | 57 | 6d | 4a | 4f | 4c | **0a** | | | | | |
| **10** | Device Name (padded with 00) ||||||||||||||| 01 |
| **20** | 03 | $D$ | 00 | b0 | $D$ | 00 | 01 | $A$ | $D_r$ | $S_r$ | $t_r$ | 00 | rekordbox ||||
| **30** | 00 | 00 | Track || 00 | 00 | $t_2$ | $t_3$ | 00 | 00 | $t_4$ | $t_5$ | 00 | 00 | 00 | 00 |
| **40** | 00 | 00 | 00 | 00 | 00 | 00 | $t_6$ | $t_7$ | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **50** | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **60** | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 01 | 00 | $U_a$ | $S_a$ | 00 | 00 | 00 | $U_l$ |
| **70** | 00 | 00 | 00 | $S_l$ | 00 | $L$ | 00 | 00 | 01 | 00 | 00 | $P_1$ | Firmware ||||
| **80** | 00 | 00 | 00 | 00 | 00 | 00 | $Sync_n$ || 00 | $F$ | ff | $P_2$ | 00 | $Pitch_1$ |||
| **90** | $l_1$ || $BPM$ || 7f | ff | ff | ff | 00 | $Pitch_2$ ||| 00 | $P_3$ | 00 | ff |
| **a0** | Beat ||| Cue || $B_b$ | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **b0** | 00 | 00 | 00 | 00 | 00 | 00 | 10 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **c0** | 00 | $Pitch_3$ || 00 | $Pitch_4$ ||| Packet ||| $nx$ | 00 | 00 | 00 |
| **d0** | 00 | 00 | 00 | 00 | | | | | | | | | | | | |

Figure 11: CDJ status packets

track is loaded from a rekordbox collection on a laptop, $S_r$ has the value 04. $t_r$ at byte 42 may be an indicator of whether the track came from rekordbox: It seems to have the value 00 when no track is loaded or an unanalyzed track is loaded, 01 when a rekordbox track is loaded, and 05 when an audio CD is loaded.

The field *rekordbox* at bytes 44–47 contains the rekordbox database ID of the loaded track when a rekordbox track is being played. Combined with the player number and slot information, this can be used to request the track metadata as described in Section 5. When an audio CD is loaded, this just contains the track index on the disc.

The track number being played (its position within a playlist or other scrolling list of tracks, as displayed on the CDJ) can be found at bytes 50 and 51, labeled *Track*. (It may be a 4-byte value and also include bytes 48 and 49, but that would seem an unmanageable number of tracks to search through.)

There are a number of bytes, labeled $t_2$ through $t_7$, whose purpose is as yet undetermined. They are all zero when there is no track loaded, but take different values when a track is loaded. (Actually, there are many other bytes than these which behave like this, so they will probably be removed from the chart in a future version of this document.)

Byte 106, labeled $U_a$ (for "USB activity"), alternates between the values 4 and 6 when there is USB activity—it may even alternate in time with the flashing USB indicator LED on the player. Byte 107, $S_a$, is the same kind of activity indicator for the SD slot. Byte 111 ($U_l$ for "USB local") has the value 4 when there is no USB media loaded, 0 when USB is loaded, and 2 when the USB Stop button has been pressed and the USB media is being unmounted.

Byte 115 ($S_l$ for "SD local") has the value 4 when there is no SD media loaded, 0 when SD is loaded, and 2 when the SD door has been opened and the SD media is being unmounted.

Byte 117, labeled $L$ (for "Link available"), appears to have the value 1 whenever USB, SD, or CD media is present in any player on the network, whether or not the Link option is chosen in the other players, and 0 otherwise.

Byte 123, labeled $P_1$, appears to describe the current play mode. The values that have been seen so far, and their apparent meanings, are:

| 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
|---|---|---|---|---|---|---|---|
| 1 | Play | Master | Sync | On-Air | 1 | 0 | 0 |

Figure 12: CDJ state flag bits

| Value | Meaning |
|---|---|
| 0 | No track is loaded |
| 2 | A track is in the process of loading |
| 3 | Player is playing normally |
| 4 | Player is playing a loop |
| 5 | Player is paused anywhere other than the cue point |
| 6 | Player is paused at the cue point |
| 7 | Cue Play is in progress (playback while the cue button is held down) |
| 8 | Cue scratch is in progress |
| 9 | Player is searching forwards or backwards |
| 17 | Player reached the end of the track and stopped |

The *Firmware* value at bytes 124–127 is an ASCII represenation of the firmware version running in the player.

The value $Sync_n$ at bytes 134–135 seems to increment whenever the player syncs to a new tempo master (another player or the mixer). I am assuming it is just a 2-byte value, because I tried syncing 256 times and saw the counter expand from byte 135 to include byte 134. It may actually be a 4-byte value and also involve bytes 132 and 133, but I wasn't going to try changing sync 65,536 times or more to find out. Perhaps we could write software to test that someday by forcing tempo master changes.

Byte 137, labeled $F$, is a bit field containing some very useful state flags, detailed in Figure 12.[9] It seems to only be available on nexus players, and others always send 0 for this byte?

Byte 139, labeled $P_2$ seems to be another play state indicator, having the value 122 (`0x7a`) when playing and 126 (`0x7e`) when stopped. When the CDJ is trying to play, but is being held in place by the DJ holding down on the jog wheel, $P_1$ considers it to be playing (value 3), while $P_2$ considers it to be stopped (value 126). Non-nexus players seem to use the value 106 (`0x6a`) when playing and 110 (`0x6e`) when stopped, while nxs2 players use the values 250 and 254 (`0xfa` and `0xfe`) so this seems to be another bit field like $F$.

There are four different places where pitch information appears in these packets: $Pitch_1$ at bytes 141–143, $Pitch_2$ at bytes 153–155,

---

[9]We have not yet seen any other values for bits 0, 1, 2, or 7 in $F$, so we're unsure if they also carry meaning. If you ever find different values for them, please let us know by filing an Issue at `https://github.com/brunchboy/dysentery/issues`

$Pitch_3$ at bytes 193–195, and $Pitch_4$ at bytes 197–199.

Each of these values represents a three-byte pitch adjustment percentage, where `0x100000` represents no adjustment (0%), `0x000000` represents slowing all the way to a complete stop (−100%, reachable only in Wide tempo mode), and `0x200000` represents playing at double speed (+100%).

Note that if playback is stopped by pushing the pitch fader all the way to −100% in Wide mode, both $P_1$ and $P_2$ still show it as playing, which is different than when the jog wheel is held down, since $P_2$ shows a stop in the latter situation.

Here is how the pitch adjustment percentage represented by $Pitch_1$ would be calculated:

$$100 \times \frac{(byte[141] \times 65536 + byte[142] \times 256 + byte[143]) - 1048576}{1048576}$$

We don't know why there are so many copies of the pitch information, or all circumstances under which they might differ from each other, but it seems that $Pitch_1$ and $Pitch_3$ report the current pitch adjustment actually in effect (as reflected on the BPM display), whether it is due to the local pitch fader, or a synced tempo master.

$Pitch_2$ and $Pitch_4$ are always tied to the position of the local pitch fader, unless Tempo Reset is active, effectively locking the pitch fader to 0% and $Pitch_2$ and $Pitch_4$ to `0x100000`, or the player is paused or the jog wheel is being held down, freezing playback and locking the local pitch to −100%, in which case they both have the value `0x000000`.

When playback stops, either due to the play button being pressed or the jog wheel held down, the value of $Pitch_4$ drops to `0x000000` instantly, while the value of $Pitch_2$ drops over time, reflecting the gradual slowdown of playback which is controlled by the player's brake speed setting. When playback starts, again either due to the play button being pressed or the jog wheel being released, both $Pitch_2$ and $Pitch_4$ gradually rise to the target pitch, at a speed controlled by the player's release speed setting.

If the player is *not* synced, but the current pitch is different than what the pitch fader would indicate (in other words, the player is in the mode where it tells you to move the pitch fader to the current BPM in order to change the pitch), moving the pitch fader changes the values of $Pitch_2$ and $Pitch_4$ until they match $Pitch_1$ and $Pitch_3$ and begin to affect the actual effective pitch. From that point on, moving the pitch fader sets the value of all of $Pitch_1$, $Pitch_2$, $Pitch_3$, and $Pitch_4$. This all seems more complicated than it really needs to be...

The current BPM of the track (the BPM at the point that is currently being played, or at the location where the player is currently

paused) can be found at bytes 146–147 (labeled $BPM$). It is a two-byte integer representing one hundred times the current track BPM. So, the current track BPM value to two decimal places can be calculated as:

$$\frac{byte[146] \times 256 + byte[147]}{100}$$

In order to obtain the actual playing BPM (the value shown in the BPM display), this value must be multiplied by the current effective pitch, calculated from $Pitch_1$ as described above. Since calculating the effective BPM reported by a CDJ is a common operation, here a simplified equation that results in the effective BPM to two decimal places, by combinining the $BPM$ and $Pitch1$ values:

$$\frac{(b[146] \times 256 + b[147]) \times (b[141] \times 65536 + b[142] \times 256 + b[143])}{104857600}$$

Because Rekordbox and the CDJs support tracks with variable BPM, this value can and does change over the course of playing such tracks. When no track is loaded, $BPM$ has the value `0xffff`.

The meaning of value $l_1$ (bytes 144–145) is not currently known. It may simply reflect whether a track is loaded or not: it seems to have the value `0x7fff` when no track is loaded, `0x8000` when a rekordbox track is loaded, and `0x0000` when a non-rekordbox track (like from a physical CD) is loaded.

Byte 157 (labeled $P_3$) seems to communicate additional information about the current play mode, with the following meanings that we have found so far:

| Value | Meaning |
|------:|---------|
| 0 | No track is loaded |
| 1 | Player is paused or playing in Reverse mode |
| 9 | Player is playing in Forward mode with jog mode set to Vinyl |
| 13 | Player is playing in Forward mode with jog mode set to CDJ |

The 4-byte beat counter (which counts each beat from 1 through the end of the track) is found in bytes 160–163, labeled $Beat$. When the player is paused at the start of the track, this seems to hold the value 0, even though it is beat 1, and when no rekordbox-analyzed track is loaded, *and in packets from non-nexus players*, this holds the value `0xffffffff`.

The counter $B_b$ at byte 166 counts out the beat within each bar, cycling $1 \rightarrow 2 \rightarrow 3 \rightarrow 4$ repeatedly, and can be used to identify the down beat (as is used in the Master Player display on the CDJs as a mixing aid). Again, when no reckordbox-analyzed track is loaded, this

holds the value 0. If you want to synchronize events to the down beat, use the CDJ status packets' $F$ value to identify the master player, but use the beat packets sent by that player (described in Section 3) to determine when the beats are actually happening.

A countdown timer to the next saved cue point is available in bytes 164–165 (labeled *Cue*). If there is no saved cue point after the current play location in the track, or if it is further than 64 bars ahead, these bytes contain the value `0x01ff` and the CDJ displays "--.- bars". As soon as there are just 64 bars (256 beats) to go before the next cue point, this value becomes `0x0100`. This is the point at which the CDJ starts to display a countdown, which it displays as "63.4 bars". As each beat goes by, this value decreases by 1, until the cue point is about to be reached, at which point the value is `0x0001` and the CDJ displays "00.1 bars". On the beat on which the cue point was saved the value is `0x0000` and the CDJ displays "00.0 Bars". On the next beat, the value becomes determined by the next cue point (if any) in the track.

Bytes 200-203 seem to contain a 4-byte packet counter labeled *Packet*, which is incremented for each packet sent by the player. (I am just guessing it is four bytes long, I have not yet watched long enough for the count to need more than the last three bytes).

Byte 204, labeled $nx$, seems to have the value `0x0f` for nexus players, and `0x05` for older players.

## 4.3   Rekordbox Status Packets

Rekordbox sends status packets which appear to be essentially identical to those sent by a mixer, as shown in Figure 10, sending "rekordbox" as its device name. The device number $D$ (bytes 33 and 36) seems to be 41 (`0x29`), although it will probably use conflict resolution to pick an unused number if multiple copies are running. The $F$ value we have seen remains consistent as a status flag, showing `0xc0` which would indicate that it is always "playing" but not synced, tempo master, nor on the air. The $BPM$ value seems to track that of the master player, and the same potential pitch values (fixed at `0x100000`, or +0%) are present, as is $X$. $B_b$ always seems to be zero.

# 5   Track Metadata

Thanks to @EvanPurkhiser[10], we finally started making progress in re-trieving metadata from CDJs, and now some shared code from Austin Wright[11] is boosting our understanding considerably!

---

[10]`https://github.com/EvanPurkhiser`
[11]`https://bitbucket.org/awwright/libpdjl`

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| **00** | 00 | 00 | 00 | 0f | R | e | m | o | t | e | D | B | S | e | r | v |
| **10** | e | r | 00 | | | | | | | | | | | | | |

Figure 13: DB Server query packet

To be polite about it, the first step is to determine the port on which the player is offering its remote database server. That can be determined by opening a TCP connection to port $12,523$ on the player and sending it sending a packet with the content shown in Figure 13.

The player will send back a two-byte response, containing the high byte of the port number followed by the low byte. So far, the response from a CDJ has always indicated a port number of 1051, but using this query to determine the port to use will protect you against any future changes. The same query can also be sent to a laptop running rekordbox to find the rekordbox database server port, which can also be queried for metadata in the exact same way described below.

To find the metadata associated with a particular track, given its rekordbox ID number, as well as the player and slot from which it was loaded (all of which can be determined from a CDJ status packet received by a virtual CDJ as described in Section 4), open a TCP connection to the device from which the track was loaded, using the port that it gave you in response to the DB Server query packet, then send the following four packets.

The first packet sent to the database server contains the five bytes `11 00 00 00 01`, and results in the same five bytes being sent back.

All further packets have a shared structure. They consist of lists of type-tagged fields (a type byte, followed some number of value bytes, although in the case of the variable-size types, the first four bytes are a big-endian integer that specifies the length of the additional value bytes that make up the field). So far, there are four known field types, and it turns out that the packet we just saw is one of them, it represents the number 1 as a 4-byte integer.

## 5.1   Field Types

The first byte of a field identifies what type of field is coming. The values `0x0f`, `0x10`, and `0x11` are followed by 1, 2, and 4 byte fixed-length integer fields, while `0x14` and `0x26` introduce variable-length fields, a binary blob and a UTF-16 big-endian string respectively.

Figure 14: Number Fields of length 1, 2, and 4



Figure 15: Binary (Blob) Field

### 5.1.1  Number Fields

Number fields are indicated by an initial byte 0x0f, 0x10, or 0x11 which is followed by big-endian integer value of length 1, 2, or 4 bytes respectively, as shown in Figure 14. So, as noted above, the initial greeting packet sent to and received back from the database server is a number field, four bytes long, representing the value 1.

### 5.1.2  Binary Fields

Variable-length binary (blob) fields are indicated by an initial byte 0x14, followed by a 4 byte big-endian integer which specifies the length of the field payload. The length is followed by the specified number of bytes (for example, an album art image, waveform or beat grid). This is illustrated in Figure 15.

### 5.1.3  String Fields

Variable-length string fields are indicated by an initial byte 0x26, followed by a 4 byte big-endian integer which specifies the length of the

Figure 16: String Field

string, in two-byte UTF-16 big-endian characters. So the length is folowed by $2 \times length$ bytes containing the actual string characters. The last character of the string is always NUL, represented by `0x0000`. This is illustrated in Figure 16.

## 5.2  Messages

Messages are introduced by a 4 byte Number field containing the magic value `0x872349ae`. This is followed by another 4 byte number field that contains a transaction ID, which starts at 1 and is incremented for each query sent, and all messages sent in response to that query will contain the same transaction ID. This is followed by a 2 byte number field containing the message type, a 1 byte number field containing the number of argument fields present in the message, and a blob field containing a series of bytes which identify the types of each argument field. This blob is always 12 bytes long, regardless of how few arguments there are (and presumably this means no message ever has more than 12 arguments). Tag bytes past the actual argument count of the message are set to 0.

The argument type tags use different values than the field type tags themselves, for some reason, and it is not clear why this redundant information is necessary at all, but that is true a number of places in the protocol as you will see later. Here are the known tag values and their meanings:

| Tag | Meaning |
| --- | --- |
| `0x02` | A string in UTF-16 big-endian encoding, with trailing NUL (zero) character |
| `0x03` | A binary blob |
| `0x06` | A 4 byte big-endian integer |

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

| 00 | 11 | 872349ae | 11 | $TxID$ | 10 | $type$ | 0f | $n$ | 14 |
| 10 | 0000000c (12) | $t_1$ $t_2$ $t_3$ $t_4$ $t_5$ $t_6$ $t_7$ $t_8$ $t_9$ $t_{10}$ $t_{11}$ $t_{12}$ |

Arguments

Figure 17: Message Header

I am guessing that if we ever see them, a tag of 4 would represent a 1 byte integer, and 5 would represent a 2 byte integer. But so far no such messages have been seen.

This header is followed by the fields that make up the message arguments, if any. The header structure is illustrated in Figure 17, where $TxID$ is the transaction ID, $n$ is the number of arguments found in the message, and $t_1$ through $t_{12}$ are the type tags for each argument, or 0 if there is no argument in that position.

Before you can send your first actual query, you need to send a special message which seems to be necessary for establishing the context for queries. It has a $type$ of 0, a special $TxID$ value of 0xfffffffe, and a single numeric argument, as shown in Figure 18.

The value $D$ is, like in the other packets we have seen, a player device number. In this case it is the device that is asking for metadata information. It must be a valid player number between 1 and 4, and that player must actually be present on the network, must not be the same player that you are contacting to request metadata from, and must not be a player that has connected to that player via Link and loaded a track from it. So the safest device number to use is the device number you are using for your virtual CDJ, but since it must be between 1 and 4, you can only do that if there are fewer than four actual CDJs on the network.

The player will respond with a message of type 0x4000, which is the common "success" response when requested data is available. The response message has two numeric arguments, the first of which is the message type of the request we sent (which was 0), and the second usually tells you the number of items that are available in response to the query you made, but in this special setup query, it returns its own

|      | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|------|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
|      |   | *start* | | | | | | *TxID* | | | | | *type* | | *args* | *tags* |
| **00** | 11 | 872349ae | | | | 11 | fffffffe | | | | 10 | 0000 | | 0f | 01 | 14 |
| **10** | 0000000c (12) | | | | 06 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | $D_{ours}$ | | | | | | | | | | | | | | |

Figure 18: Query context setup message

|      | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|------|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
|      |   | *start* | | | | | | *TxID* | | | | | *type* | | *args* | *tags* |
| **00** | 11 | 872349ae | | | | 11 | fffffffe | | | | 10 | 4000 | | 0f | 02 | 14 |
| **10** | 0000000c (12) | | | | 06 | 06 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | 00000000 | | | | 11 | $D_{theirs}$ | | | | | | | | | |

Figure 19: Query context setup response

player number. The overall structure is illustrated in Figure 19.

## 5.3   Track Metadata

To ask for metadata about a particular track, send a packet like the one shown in Figure 20.

As described above, $TxID$ should be 1 for the first query packet you send, 2 for the next, and so on. $D$ should have the same value you used in your initial query context setup packet, identifying the device that is asking the question. $S_r$ is the slot in which the track

|      | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|------|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
|      |   | *start* | | | | | | *TxID* | | | | | *type* | | *args* | *tags* |
| **00** | 11 | 872349ae | | | | 11 | $TxID$ | | | | 10 | 2002 | | 0f | 02 | 14 |
| **10** | 0000000c (12) | | | | 06 | 06 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | $D$ | 01 | $S_r$ | 01 | 11 | *rekordbox* | | | | | | | | | |

Figure 20: Track metadata request message

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | start | | | | | *TxID* | | | | *type* | | *args* | | *tags* |

|  | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| **00** | 11 | 872349ae | | | | 11 | *TxID* | | | 10 | 4000 | | 0f | 02 | 14 |
| **10** | 0000000c (12) | | | | 06 | 06 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | 00002002 | | | 11 | 0000000b | | | | | | | | | |

Figure 21: Track metadata ready response

being asked about can be found, and has the same values used in CDJ status packets, as shown in Figure 11. Similarly, *rekordbox* identifies the local rekordbox database ID of the track being asked about, as found in the CDJ status packet.

Track metadata requests are built on the mechanism that is used to request and draw scrollable menus on the CDJs, so the request is essentially interpreted as setting up to draw the "menu" of information that is known about the track. The player responds with a success indicator, saying it is ready to send these "menu items" and reporting how many of them are available, as shown in Figure 21.

We've seen this type of "data available" response already in Figure 19, but this one is a more typical example. As usual, $TxID$ matches the value we sent in our request, and the first argument, with value 0x2002, reflects the *type* field of our request. The second argument reports that there are 11 (0x0b) entries of track metadata available to be retrieved for the track we asked about, and that the player is ready to send them to us.

If there was no track with ID *rekordbox* in that media slot, the second argument would have the value 0xffffffff to let us know. If we messed up something else about the request, we will get a response with a *type* other than 0x4000. See Section 5.12 for instructions on how to explore these variations on your own.

But assuming everything went well, we can get the player to send us all eleven of those metadata entries by telling it to render all of the current menu, using a "render menu" request with *type* 0x3000 shown in Figure 22.

As always, the value of $TxID$ should be one higher than the one you sent in your setup packet, while the values of $D$ and $S_r$ should be identical to what you sent in it.

The request has six numeric arguments. At this point it is worth talking a bit more about the byte after $D$ in the first argument. This seems to specify the location of the menu being drawn, with the value 1 meaning the main menu on the left-hand half of the CDJ display,

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | start | | | | | TxID | | | | | type | | args | | tags |
| 00 | 11 | 872349ae | | | | 11 | $TxID$ | | | | 10 | 3000 | | 0f | 06 | 14 |
| 10 | 0000000c (12) | | | | 06 | 06 | 06 | 06 | 06 | 06 | 00 | 00 | 00 | 00 | 00 | 00 |
| 20 | 11 | $D$ | 01 | $S_r$ | 01 | 11 | $offset$ | | | | 11 | $limit$ | | | | 11 |
| 30 | 00000000 | | | | 11 | $limit_2$ | | | | 11 | 00000000 | | | | | |

Figure 22: Render Menu request message

while 2 means the submenu (for example the info popup when it is open) which overlays the right-hand half of the display. We don't yet know exactly what, if any, difference there is between the response details when 2 is used instead of 1 here. And special data requests use different values: for example, the track waveform summary, which is drawn in a strip along the entire bottom of the display, is requested with a menu location number of 8 in this second byte.

The byte after $S_r$ seems to always have the value 1. We've seen it described as "source analyzed" but don't know exactly what that means.

The second argument, $offset$, specifies which menu entry is the first one you want to see, and the third argument, $limit$, specifies how many should be sent. In this case, since there are only 11 entries, we can request them all at once, so we will set $offset$ to 0 and $limit$ to 11. But for large playlists, for example, you need to request batches of entries that are smaller than the total available, or the player will be unable to send them to you. We have not found what the exact limit is, but getting 64 at a time seems to work on Nexus 2 players.

We don't know the purpose of the fourth argument, but sending a value of 0 works. The fifth argument, $limit_2$, is also poorly understood, but sending a second copy of $limit$ here works. And the sixth and final argument also has an unknown purpose, but 0 works.

So, for our metadata request, the packet we want to send in order to get all the metadata will have the specific values shown in Figure 23:

This causes the player to send us 13 messages: The 11 metadata items we requested are sent (with $type$ 0x4101, Figure 25), but they are preceded by a menu header message (with $type$ 0x4001, Figure 24), and followed by a menu footer message (with $type$ 0x4201, Figure 26). This wrapping happens with all "render menu" requests, and the menu footer is an easy way to know you are done, although you can also count the messages.

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | | start | | | | | TxID | | | | type | | args | | tags |
| 00 | 11 | | 872349ae | | | 11 | | TxID | | | 10 | | 3000 | 0f | 06 | 14 |
| 10 | 0000000c (12) | | | | 06 | 06 | 06 | 06 | 06 | 06 | 00 | 00 | 00 | 00 | 00 | 00 |
| 20 | 11 | $D$ | 01 | $S_r$ | 01 | 11 | | 00000000 | | | 11 | | 0000000b | | | 11 |
| 30 | | 00000000 | | | 11 | | 0000000b | | 11 | | 00000000 | | | | | |

Figure 23: Render track metadata request message

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | | start | | | | | TxID | | | | type | | args | | tags |
| 00 | 11 | | 872349ae | | | 11 | | TxID | | | 10 | | 4001 | 0f | 02 | 14 |
| 10 | 0000000c (12) | | | | 06 | 06 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| 20 | 11 | | 00000001 | | | 11 | | 00000000 | | | | | | | | |

Figure 24: Menu header response

The menu item responses all have the same structure, and use all twelve message argument slots, containing ten numbers and two strings, although they generally don't have meaningful values in all of the slots. They have the general structure shown in in Figure 25, and the arguments are:

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

|  | start |  |  |  |  |  | *TxID* |  |  |  | *type* |  | *args* |  | *tags* |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

| **00** | 11 | 872349ae | 11 | *TxID* | 10 | 4101 | 0f | 0c | 14 |
|---|---|---|---|---|---|---|---|---|---|

| **10** | 0000000c (12) | 06 | 06 | 06 | 02 | 06 | 02 | 06 | 06 | 06 | 06 | 06 | 06 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

Arguments

Figure 25: Menu item response

| Arg | Type | Meaning |
|---|---|---|
| 1 | number | parent ID, such as an artist for a track item |
| 2 | number | main ID, such as *rekordbox* for a track item |
| 3 | number | length in bytes of Label 1 |
| 4 | string | Label 1 (main text, such as the track title or artist name, as appropriate for the item type) |
| 5 | number | length in bytes of Label 2 |
| 6 | string | Label 2 (secondary text, e.g. artist name for playlist entries, where Label 1 holds the title) |
| 7 | number | type of this item (see Section 5.5) |
| 8 | number | some sort of flags field, details still unclear |
| 9 | number | holds *artwork* ID when type is Track Title |
| 10 | number | playlist position when relevant, e.g. when listing a playlist |
| 11 | number | unknown |
| 12 | number | unknown |

### 5.3.1   Track Metadata Item 1: Title

The first item returned after the menu header is the track title, so argument 7 has the value 0x04. Argument 1, which may always be some kind of parent ID, holds the artist ID associated with the track. The second argument seems to always be the main ID, and for this response it holds the *rekordbox* ID of the track. Argument 4 holds the track title text, and argument 8 holds the album *artwork* ID if any is available for the track. This ID can be used to retrieve the actual

album art image as described in Section 5.6.

### 5.3.2   Track Metadata Item 2: Artist

The second item contains artist information so argument 7 has the value `0x07`. Argument 2 holds the artist ID, argument 4 contains the text of the artist name.

### 5.3.3   Track Metadata Item 3: Album Title

The third item contains album title information so argument 7 has the value `0x02`. Argument 4 contains the text of the album name.

### 5.3.4   Track Metadata Item 4: Duration

The fourth item contains track duration information so argument 7 has the value `0x0b`. Argument 2 contains the length, in seconds, of the track when played at normal tempo.

### 5.3.5   Track Metadata Item 5: Tempo

The fifth item contains tempo information so argument 7 has the value `0x0d`. Argument 2 contains the track's starting tempo, in BPM, times 100, as reported in $BPM$ values in other packets described earlier.

### 5.3.6   Track Metadata Item 6: Comment

The sixth item contains comment information so argument 7 has the value `0x23`. Argument 4 contains the text of the track comment entered by the DJ in rekordbox.

### 5.3.7   Track Metadata Item 7: Key

The seventh item contains key information so argument 7 has the value `0x0f`. Argument 4 contains the text of the track's dominant key signature.

### 5.3.8   Track Metadata Item 8: Rating

The eighth item contains rating information so argument 7 has the value `0x0a`. Argument 2 contains a value from 0 to 5 corresponding to the number of stars the DJ has assigned the track in rekordbox.

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | | *start* | | | | | *TxID* | | | | *type* | | *args* | | *tags* |
| **00** | 11 | | 872349ae | | | 11 | | *TxID* | | | 10 | 4201 | | 0f | 00 | 14 |
| **10** | 0000000c (12) | | | | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |

Figure 26: Menu footer response

### 5.3.9    Track Metadata Item 9: Color

The ninth item contains color information so argument 7 has the value `0x13`. I assume Argument 2 contains a value corresponding to the color the DJ has assigned the track in rekordbox, but I have not yet figured out the details.

### 5.3.10    Track Metadata Item 10: Genre

The tenth item contains genre information so argument 7 has the value `0x06`. Argument 2 contains the numeric genre ID, and argument 4 contains the text of the genre name.

### 5.3.11    Track Metadata Item 11: Date Added

The eleventh and final item contains the date added information so argument 7 has the value `0x2e`. Argument 4 contains the date the track was added to the collection in the format "yyyy-mm-dd". This information seems to propagate into rekordbox from iTunes.

## 5.4    Menu Footer Response

The menu footer message has a *type* of `0x4201` and no arguments, so it has a header only, and is always made up of the exact same sequence of bytes (apart from the *TxID*), as shown in Figure 26.

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | | start | | | | | *TxID* | | | | *type* | | *args* | | *tags* |
| **00** | 11 | | 872349ae | | | 11 | | *TxID* | | | 10 | 2003 | | 0f | 02 | 14 |
| **10** | 0000000c (12) | | | | 06 | 06 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | $D$ | 08 | $S_r$ | 01 | 11 | | *artwork* | | | | | | | | |

Figure 27: Track artwork request message

## 5.5  Menu Item Types

As noted above, the seventh argument in a menu item response identifies the type of the item. The meanings we have identified so far are:

| Type | Meaning |
|---|---|
| 0x0001 | Folder (such as in the playlists menu)[12] |
| 0x0002 | Album title |
| 0x0003 | Disc |
| 0x0004 | Track Title |
| 0x0006 | Genre |
| 0x0007 | Artist |
| 0x0008 | Playlist |
| 0x000a | Rating |
| 0x000b | Duration (in seconds) |
| 0x000d | Tempo |
| 0x000f | Key |
| 0x0013 | Color |
| 0x0023 | Comment |
| 0x002e | Date Added |
| 0x0704 | Track List |

As noted above, track metadata responses use many of these types. Others are used in different kinds of menus and queries.

## 5.6  Album Art

To request the artwork image associated with a track, send a message with *type* 0x2003 containing the *artwork* ID that was specified in the track title item (as described in Section 5.3.1), like the one shown in Figure 27.

---

[12]A nested list of playlists rather than an individual playlist.

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | *start* | | | | | | *TxID* | | | *type* | | | *args* | | *tags* |
| **00** | 11 | \| 872349ae | | | \| 11 | \| *TxID* | | | \| 10 | \| 4002 | | \| 0f | \| 04 | \| 14 |
| **10** | 0000000c (12) | \| 06 | 06 | 06 | 03 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 \| 00002003 | | \| 11 \| 00000000 | | \| 11 \| *length* | | \| 14 |

| | | |
|---|---|---|
| *length* | Image bytes | |

Figure 28: Track artwork response message

As usual, *seq* should be incremented each time you send a query, and will be used to identify the response messages. $D$ should have the same value you used in your initial query context setup packet, identifying the device that is asking the question. $S_r$ is the slot in which the track being asked about can be found, and has the same values used in CDJ status packets, as shown in Figure 11. Finally, *artwork* identifies the specific artwork image you are requesting, as it was specified in the track medatata response. As with other graphical requests, the value after $D$, which identifies the location of the menu for which data is being loaded, is 8.

The response will be a message with *type* 0x4002, containing four arguments. The first argument echoes back our request type, which was 0x2003. The second always seems to be 0. The third contains the length of the image in bytes, which seems redundant, because that is also conveyed in the fourth argument itself, which is a blob containing the actual bytes of the image data, as shown in Figure 28.

To experiment with this, start up dysentery in a Clojure REPL and connect to a player as described in Section 5.12, then evaluate an expression like:

```
(def art (db/request-album-art p2 3 3))
```

Replace the arguments with the var holding your player connection, the proper $S_r$ number for the media slot the art is found in, and the *artwork* ID of the album art, and dysentery will open a window like Figure 29 showing the image.

Figure 29: Example album art window

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | | *start* | | | | | | *TxID* | | | *type* | | | *args* | *tags* |
| **00** | 11 | | 872349ae | | | 11 | | | *TxID* | | 10 | | 2204 | 0f | 02 | 14 |
| **10** | 0000000c (12) | | | | 06 | 06 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | $D$ | 08 | $S_r$ | 01 | 11 | | *rekordbox* | | | | | | | | |

Figure 30: Track beat grid request message

## 5.7   Beat Grids

The CDJs do not send any timing information other than beat numbers during playback, which has made it difficult to offer absolute timecode information. The discovery of beat grid requests provides a clean answer to the problem. The beat grid for a track is a list of every beat which occurs in the track, along with the point in time at which that beat would occur if the track were played at its standard (100%) tempo. Armed with this table, it is possible to translate any beat packet into an absolute position within the track, and, combined with the tempo information, to generate timecode signals allowing other software (such as video resources) to sync tightly with DJ playback.

To request the beat grid of a track, send a message with *type* 0x2204 containing the *rekordbox* ID of the track, like the one shown in Figure 30.

As usual, *seq* should be incremented each time you send a query, and will be used to identify the response messages. $D$ should have the same value you used in your initial query context setup packet, identifying the device that is asking the question. $S_r$ is the slot in which the track being asked about can be found, and has the same values used in CDJ status packets, as shown in Figure 11. Finally, *rekordbox* identifies the specific artwork image you are requesting, as

|     | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|-----|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
|     |   | **start** | | | | | | *TxID* | | | | *type* | | *args* | | *tags* |
| **00** | 11 | 872349ae | | | | 11 | *TxID* | | | | 10 | 4602 | | 0f | 04 | 14 |
| **10** | 0000000c (12) | | | | | 06 | 06 | 06 | 03 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | 00002204 | | | | 11 | 00000000 | | | | 11 | *length* | | | | 14 |

*length*          Beat grid bytes

Figure 31: Track beat grid response message

found in a CDJ status packet or playlist response. As with graphical requests, the value after $D$, which identifies the location of the menu for which data is being loaded, is 8.

The response will be a message with *type* 0x4602, containing four arguments. The first argument echoes back our request type, which was 0x2204. The second always seems to be 0. The third contains the length of the beat grid in bytes, which seems redundant, because that is also conveyed in the fourth argument itself, which is a blob containing the actual bytes of the beat grid, as shown in Figure 31.

The beat grid itself is spread through the value returned as argument 4, consisting of one-byte beat-within-bar numbers (labeled $B_b$ in Figure 11), followed by four-byte timing information, specifying the number of milliseconds after the start of the track (when played at its native tempo) at which that beat falls.

The $B_b$ value of the first beat in the track is found at byte 20 of argument 4, and the time at which that beat occurs, in milliseconds, is encoded as a 4-byte big-endian integer at bytes 21–24. Subsequent beats are found at 16-byte intervals, so the second $B_b$ value is found at byte 36, and the second beat's time, in milliseconds from the start of the track, is the big-endian integer at bytes 37–40. The $B_b$ value for the third beat is at byte 52, and its millisecond time is at bytes 53–56, and so on.

The purpose of the other bytes within the beat grid is so far undetermined. It looks like there may be some sort of monotonically in-

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | | *start* | | | | | *TxID* | | | | *type* | | *args* | | *tags* |
| **00** | 11 | | 872349ae | | | 11 | | $TxID$ | | | 10 | 2004 | | 0f | 05 | 14 |
| **10** | 0000000c (12) | | | | | 06 | 06 | 06 | 06 | 03 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | $D$ | 08 | $S_r$ | 01 | 11 | 00000004 | | | | 11 | *rekordbox* | | | | 11 |
| **30** | 00000000 | | | | | | | | | | | | | | | |

Figure 32: Waveform preview request message

creasing value following the beat millisecond value, but what it means, and why it sometimes skips values, is mysterious.

## 5.8   Requesting Track Waveforms

Waveform data for tracks can be requested, both the preview, which is 300 pixels long, and the detailed waveform, which uses 150 pixels per second of track content. Although Austin has figured out how to interpret the bytes for non-color waveforms, all of my players seem to return an expanded format that we have not yet figured out how to interpret and draw properly, so this section is incomplete.

To request the waveform preview of a track, send a message with *type* 0x2004 containing the *rekordbox* ID of the track, like the one shown in Figure 32.

As usual, *seq* should be incremented each time you send a query, and will be used to identify the response messages. $D$ should have the same value you used in your initial query context setup packet, identifying the device that is asking the question. $S_r$ is the slot in which the track being asked about can be found, and has the same values used in CDJ status packets, as shown in Figure 11. Finally, *rekordbox* identifies the specific artwork image you are requesting, as found in a CDJ status packet or playlist response. As with graphical requests, the value after $D$, which identifies the location of the menu for which data is being loaded, is 8.

You may have noticed that the argument list of the message in Figure 32 specifies that there are five arguments, but in fact the message contains only the first four, numeric, arguments. The fifth, blob, argument is missing. This seems to imply the blob is empty, and this very strange feature of the protocol is, in fact, the way the track metadata is requested. The fifth argument must be specified in the message header but not actually present.

The second argument has an unknown purpose, but we have seen

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | | start | | | | | *TxID* | | | | *type* | | *args* | | *tags* |
| **00** | 11 | | 872349ae | | | 11 | | *TxID* | | | 10 | 4402 | | 0f | 04 | 14 |
| **10** | 0000000c (12) | | | | | 06 | 06 | 06 | 03 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | | 00002004 | | | 11 | | 00000000 | | | 11 | | *length* | | | 14 |

| *length* | Waveform preview bytes |
|---|---|

Figure 33: Waveform preview response message

values of 3 or 4 for it. Similarly, the fourth argument's meaning is unknown, but it always seems to be 0.

The response will be a message with *type* 0x4402, containing five arguments. The first argument echoes back our request type, which was 0x2004. The second always seems to be 0. The third contains the length of the waveform preview in bytes, which seems redundant, because that is also conveyed in the fourth argument itself, which is a blob containing the actual bytes of the waveform preview, as shown in Figure 33.

Austin reported that 300-byte values were being returned in the preview blob, each representing a pixel-wide slice of waveform, and the way to interpret them was to mask off the 5 low-order bits of each value, which represented the height of the waveform at that point, from 0 to 31. The 3 high-order bits represented the saturation of the waveform, ranging from white at 111 to blue at 000. This may have been the behavior of earlier firmware versions, or different player models? The results I have obtained with the CDJ-2000 nexus are different.

In my testing, there are 900 bytes of waveform data returned. The first 800 of them contain 400 columns of waveform data, in the form of two-byte pairs, where the first byte is the height of the waveform at that column (a value ranging from 0 to 31), and the second byte is the whiteness, as before, where 0 is blue and 7 is fully white. My players seem to only pay attention to the highest bit of whiteness, drawing the waveform as either very dark or light blue accordingly.

Figure 34: Example waveform preview window



Figure 35: Full track list request message

To experiment with this, start up dysentery in a Clojure REPL and connect to a player as described in Section 5.12, then evaluate an expression like (db/request-waveform-preview p2 3 1060), replacing the arguments with the var holding your player connection, the proper $S_r$ number for the media slot the track is found in, and the *rekordbox* ID of the track, and dysentery will open a window like Figure 34 showing the waveform preview.

I don't yet know what the remaining 100 bytes mean; perhaps they are color information for players that support colored waveforms?

Requesting the detailed waveform is very similar to requesting the preview, but the request type and arguments are slightly different. The details will be documented once it is implemented in dysentery.

## 5.9   Requesting Cue Points

The locations of the cue points stored in a track can be obtained with another request. The details will be documented once it is implemented in dysentery.

## 5.10   Requesting All Tracks

If you want to cache all the metadata associated with a media stick, this query is a good starting point. Send a packet like the one shown in Figure 35.

As always, $TxID$ should be 1 for the first query packet you send,

2 for the next, and so on. $D$ should have the same value you used in your initial query context setup packet, identifying the device that is asking the question. $S_r$ is the slot in which the track being asked about can be found, and has the same values used in CDJ status packets, as shown in Figure 11. The new *sort* parameter determines the order in which the tracks are sorted, and that also affects the item type returned, along with the secondary information (beyond the title) that it contains about the track, as described in Section 5.10.1. We will start out assuming the tracks are being requested in title order, which can be done by sending a *sort* argument value of 0 or 1.

Track list requests (just like metadata requests) are built on the mechanism that is used to request and draw scrollable menus on the CDJs. The player responds with a success indicator, saying it is ready to send these "menu items" and reporting how many of them are available, much like shown in Figure 21, although the first argument will be `0x1004` to reflect the message type we just sent, rather than `0x2002` as it was for the metadata request.

As with metadata, the next step is to send a "render menu" request like that in Figure 22 to get the actual results. But the number of results available is likely to be much higher than shown in Figure 21, because we have asked about all tracks in the media slot. That means we will probably need more than one "render menu" request to get them all.

I don't know how many items you can safely ask for at one time. I have had success with values as high as 64 on my CDJ-2000 nexus players, but they failed when asking for numbers in the thousands. So to be safe, you should ask for the results in chunks of 64 or smaller, by setting *limit* and *limit2* to the smaller of 64 and the remaining number of results you want, and incrementing *offset* by 64 in each request until you have retrieved them all.

As with metadata requests, you will get back two more messages than you ask for, because you first get a menu header message (with *type* `0x4001`, Figure 24), then the requested menu items are sent (with *type* `0x4101`, Figure 25), and finally these are followed by a menu footer message (with *type* `0x4201`, Figure 26). This wrapping happens with all "render menu" requests, and the menu footer is an easy way to know you are done, although you can also count the messages.

The details of the menu items are slightly different than in the case of a metadata request. In the example where you are retrieving tracks in title order, they will have content like this:

| Arg | Type | Meaning |
|---|---|---|
| 1 | number | artist id |
| 2 | number | *rekordbox* id of track |
| 3 | number | length in bytes of Label 1 |
| 4 | string | Label 1, Track Title |
| 5 | number | length in bytes of Label 2 |
| 6 | string | Label 2, Artist Name |
| 7 | number | type of this item, `0x704` for this sort order |
| 8 | number | some sort of flags field, details still unclear |
| 9 | number | unknown |
| 10 | number | unknown |
| 11 | number | unknown |
| 12 | number | unknown |

### 5.10.1   Alternate Track List Sort Orders

As noted above, you can request the track list in a diferent order by supplying a different value for the *sort* parameter. The value 0 or 1 gives the order and information just described. The *sort* values discovered so far are shown in the following table, and return menu items with the specified type values in argument 7.

| Sort | Type | Description |
|---|---|---|
| 1 | `0x704` | Sorted by track title, described above |
| 2 | `0x704` | Sorted by artist name, same item type |
| 3 | `0x204` | Sorted by album, Arg 1 becomes album ID, Label 2 becomes album name |
| 4 | `0xd04` | Sorted by BPM, Arg 1 becomes BPM×100, Label 2 empty |
| 6 | `0x604` | Sorted by genre, Arg 1 becomes genre ID, Label 2 becomes genre name |
| 7 | `0x2304` | Sorted by comment, Label 2 becomes comment |

To experiment with this, start up dysentery in a Clojure REPL and connect to a player as described in Section 5.12, then evaluate an expression like `(db/request-track-list p2 3)`, replacing the arguments with the `var` holding your player connection and the proper $S_r$ number for the media slot containing the tracks you want to list. You can also add a third argument to specify a sort order, like this to sort all the tracks in the USB slot by BPM:

```
(db/request-track-list p2 3 4)
```

### 5.11   Playlists

If you want to be more selective about the metadata that you are caching, you can navigate the playlist folder hierarchy and deal with

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | a | b | c | d | e | f |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | | *start* | | | | | | *TxID* | | | | *type* | | *args* | | *tags* |
| **00** | 11 | 872349ae | | | | 11 | $TxID$ | | | | 10 | 1105 | | 0f | 04 | 14 |
| **10** | 0000000c (12) | | | | | 06 | 06 | 06 | 06 | 00 | 00 | 00 | 00 | 00 | 00 | 00 | 00 |
| **20** | 11 | $D$ | 01 | $S_r$ | 01 | 11 | *sort* | | | | 11 | *id* | | | | 11 |
| **30** | $folder?$ | | | | | | | | | | | | | | | |

Figure 36: Playlist request message

only specific playlists. This process is essentially the same as asking for all tracks, except that in the playlist request you specify the playlist or folder that you want to list. To start at the root of the playlist folder hierarchy, you request folder 0. A playlist request has the struture shown in Figure 36.

As always, $TxID$ should be 1 for the first query packet you send, 2 for the next, and so on. $D$ should have the same value you used in your initial query context setup packet, identifying the device that is asking the question. $S_r$ is the slot in which the track being asked about can be found, and has the same values used in CDJ status packets, as shown in Figure 11. You specify the ID of the playlist or folder you want to list in the $id$ argument, and set $folder?$ to 1 if you are asking for a folder, and 0 if you are asking for a playlist. As noted above, to get the top-level list of playlists, ask for folder 0, by passing an $id$ of 0 and passing $folder?$ as 1.

Much as when listing all tracks, the response may tell you there are more entries in the playlist than you can retreive in a single request, so you should follow the procedure outlined in Section 5.10 to request your results in smaller batches. The followup queries that you send are identical for playlists as they are described in that section. The actual menu items returned when you are asking for a folder have the following content:

| Arg | Type | Meaning |
|---|---|---|
| 1 | number | parent folder id |
| 2 | number | id of playlist or folder |
| 3 | number | length in bytes of Label 1 |
| 4 | string | Label 1, Name of playlist or folder |
| 5 | number | length in bytes of Label 2 |
| 6 | string | Label 2, empty |
| 7 | number | type of this item, `0x1` for folder, `0x8` for playlist |
| 8 | number | unknown |
| 9 | number | unknown |
| 10 | number | playlist position |
| 11 | number | unknown |
| 12 | number | unknown |

When you have requested a playlist (by passing its *id* and a value of 0 for *folder?*) the responses you get are track list entries, just like when you request all tracks as shown in Section 5.10. And just like in that section, you can get the results in a different order by specifying a value for *sort*. The supported values and corresponding item types returned seem to be the same as described there. Additionally, if you pass a *sort* value of 9, the playlist entries will come back sorted by track title, Label 2 will be empty, and the item type will be `0x2904`.

To experiment with this, start up dysentery in a Clojure REPL and connect to a player as described in Section 5.12, then evaluate an expression like `(db/request-playlist p2 3 1)`, replacing the arguments with the `var` holding your player connection, the proper $S_r$ number for the media slot containing the playlist you want to list, and the playlist ID. You can also add a third argument to specify that you want to list a folder, like this using folder ID 0 to request the root folder:

```
(db/request-playlist p2 3 0 true)
```

Finally, you can add a fourth argument to specify a sort order, like this to sort all the tracks in playlist 12 by genre:

```
(db/request-playlist p2 3 12 false 6)
```

## 5.12   Experimenting with Metadata

The best way to get a feel for the details of working with these messages is to load dysentery into a Clojure REPL, as described on the project page, and play with some of the functions in the `dysentery.dbserver` namespace. Have no more than three players connected and active on your network, so you have an unused player number for dysentery to use. In this example, player number 1 is available, so we set dysentery up to pose as player 1:

```
> lein repl
nREPL server started on port 53806 on host 127.0.0.1 -
 nrepl://127.0.0.1:53806
REPL-y 0.3.7, nREPL 0.2.12
Clojure 1.8.0
Java HotSpot(TM) 64-Bit Server VM 1.8.0_77-b03
dysentery loaded.
dysentery.core=> (view/watch-devices :player-number 1)
Looking for DJ Link devices...
Found:
   DJM-2000nexus 33 /172.16.42.3
   CDJ-2000nexus 2 /172.16.42.4

To communicate create a virtual CDJ with address
  /172.16.42.2, MAC address 3c:15:c2:e7:08:6c,
  and use broadcast address /172.16.42.255
:socket #object[java.net.DatagramSocket 0x22b952b1
                "java.net.DatagramSocket@22b952b1"],
 :watcher #future[:status :pending, :val nil 0x3eb8f41b]
dysentery.core> (def p2 (db/connect-to-player 2 1))
Transaction: 4294967294, message type: 0x4000
 (requested data available), argument count: 2, arguments:
  number:          0 (0x00000000) [request type]
  number:          2 (0x00000002) [# items available]
#'dysentery.core/p2
dysentery.core> (def md (db/request-metadata p2 2 1))
Sending > Transaction: 1, message type: 0x2002
 (request track metadata), argument count: 2, arguments:
  number:   16843265 (0x01010201) [player, menu, media, 1]
  number:          1 (0x00000001) [rekordbox ID]
Received > Transaction: 1, message type: 0x4000
 (requested data available), argument count: 2, arguments:
  number:       8194 (0x00002002) [request type]
  number:         11 (0x0000000b) [# items available]
Sending > Transaction: 2, message type: 0x3000
 (render menu), argument count: 6, arguments:
  number:   16843265 (0x01010201) [player, menu, media, 1]
  number:          0 (0x00000000) [offset]
  number:         11 (0x0000000b) [limit]
  number:          0 (0x00000000) [unknown (0)?]
  number:         11 (0x0000000b) [len_a (= limit)?]
  number:          0 (0x00000000) [unknown (0)?]
Received 1  > Transaction: 2, message type: 0x4001
 (rendered menu header), argument count: 2, arguments:
  number:          1 (0x00000001) [unknown]
  number:          0 (0x00000000) [unknown]
Received 2  > Transaction: 2, message type: 0x4101
 (rendered menu item), argument count: 12, arguments:
  number:          1 (0x00000001) [numeric 1 (parent id)]
```

```
  number:           1 (0x00000001) [numeric 2 (this id)]
  number:          80 (0x00000050) [label 1 byte size]
  string: "Escape ft. Zoë Phillips" [label 1]
  number:           2 (0x00000002) [label 2 byte size]
  string: "" [label 2]
  number:           4 (0x00000004) [item type: Track Title]
  number:    16777216 (0x01000000) [column configuration?]
  number:           0 (0x00000000) [album art id]
  number:           0 (0x00000000) [playlist position]
  number:         256 (0x00000100) [unknown]
  number:           0 (0x00000000) [unknown]
...
Received 13  > Transaction: 2, message type: 0x4201
 (rendered menu footer), argument count: 0, arguments:
#'dysentery.core/md
dysentery.core>
```

In this interaction, after setting up the watcher so we can find players on the network, we set the `var p2` to be a connection to player 2, in which we are posing as player 1. Then we submit a metadata request to `p2`, requesting the track in slot 2 (SD card), with *rekordbox* id 1. You can see the messages being sent and received to accomplish that. For more functions that you can call, including the very flexible `experiment` function, look at the source for the `dysentery.dbserver` namespace. Most of the response messages containing track metadata were omitted for brevity; you will get more meaningful results trying it with your own tracks, and then you can see all the details.

# 6   What's Missing?

We know this analysis isn't complete. Here are some loose ends to explore.

## 6.1   Background Research

Prior to Evan and Austin's breakthroughs, here is all we knew:

By setting up a managed switch to mirror traffic sent directly between CDJs, we have been able to see how the Link Info operation is implemented: The players open a direct TCP connection between each other, and send queries to obtain the metadata about tracks with particular rekordbox ID values.

Using an Ethernet switch with port mirroring was, as we hoped, very helpful. As can be seen in the capture at `https://github.com/brunchboy/dysentery/raw/master/doc/assets/LinkInfo.`

`pcapng`, which shows a CDJ with IP address 169.254.192.112 booting, the new CDJ opens two TCP connections to the other CDJ at 169.254.119.181.

The first session (given id 0 by Wireshark), which begins at packet 206, connecting to port 12523, determines the port to use for metadata queries.

The second TCP connection (Wireshark display filter `tcp.stream eq 1`), beginning at packet 212 and connecting to port 1051, shows the track information used by the Link Info display passing between the CDJs. You can see packets reflecting the initial display of a track that was already loaded, then new information as the linked CDJ loaded three other tracks.

There is another capture at `https://github.com/brunchboy/dysentery/raw/master/doc/assets/LinkInfo2.pcapng`, with more Link Info streams to be studied (all of the odd numbered `tcp.stream` values in Wireshark are the relevant ones).

## 6.2   Deeper Emulation and Tempo Mastery

With the help of Austin's `libpdjl`[13], it looks like it will be possible to emulate a CDJ well enough to act as a tempo master, which will allow bidirectional sync with Ableton Link. This is probably my next area of focus.

It should also be easy enough to capture the traffic that rekordbox sends when it tells a player to load and start playing a track, and reproduce that.

## 6.3   Mysterious Values

There are still many values with unknown meanings described above, and many menu types that have yet to be explored; I hav focused on the ones that will be immediately useful to Beat Link Trigger. Contributions of additional research and insight are eagerly welcomed—I would have not gotten nearly this far without help!

## 6.4   Reading Data Without a CDJ

In order to offer metadata, timecode, waveforms, and so on, when there are four actual CDJs on the network, it is necessary to pre-load and cache all the data, since dbserver queries are not possible with all player numbers in use. While this can be done with a single CDJ powered up, it would be really nice to be able to read the information right out of the rekordbox files on the media that the DJ will be using for the show. So far, we don't know how to do that.

---

[13] `https://bitbucket.org/awwright/libpdjl`

Before we discovered how to ask players for metadata about particular tracks, we did some research into the underlying rexordbox database. The database format is called DeviceSQL,[14] and there used to be a free quick start suite for working with it[15] but that site no longer exists because the original (California) company Encirq[16] was acquired by the Japanese Ubiquitous Corporation in 2008.[17] It seems to still be available,[18] but I'd be surprised if they wanted to help out an open source effort like this one.

## 6.5 CDJ Packets to Rekordbox

Performing a packet capture while rekordbox is running reveals that the CDJs send unicast packets to the rekordbox address on port 50000, in addition to the packets they normally broadcast on that port. Figuring out how to pose as rekordbox might be useful in order to see what additional data these can offer, although that may be much more work than posing as a CDJ.

## 6.6 Dysentery

If you have access to Pioneer equipment and are willing to help us validate this analysis, and perhaps even figure out more details, you can find the tool that is being used to perform this research at:
`https://github.com/brunchboy/dysentery`

# List of Figures

---

[14]`https://www.quora.com/What-database-system-did-Greg-Kemnitz-develop`
[15]`http://java.sys-con.com/node/328557`
[16]`https://www.crunchbase.com/organization/encirq-corporation`
[17]`http://www.ubiquitous.co.jp/en/news/press/pdf/p1730_01.pdf`
[18]`http://www.ubiquitous.co.jp/en/products/db/md/devicesql/`

*deep*
*Symmetry*

http://deepsymmetry.org