

# DJ Link Packet Analysis

James Elliott  
Deep Symmetry, LLC

May 2, 2016

## Abstract

The protocol used by Pioneer professional DJ equipment to communicate and coordinate performances can be monitored to provide useful information for synchronizing other software, such as light shows and sequencers. By creating a “virtual CDJ” that sends appropriate packets to the network, other devices can be induced to send packets containing even more useful information about their state. This article documents what has been learned so far about the protocol, and how to accomplish these tasks.

## 1 Mixer Startup

When the mixer starts up, after it obtains an IP address (or gives up on doing that and self-assigns an address), it sends out what look like a series of packets<sup>1</sup> simply announcing its existence to UDP port 50000 on the broadcast address of the local network.

These have a data length<sup>2</sup> of 37 bytes, appear roughly every 300 milliseconds, and have the content shown in Figure 1.




Figure 1: Initial announcement packets from Mixer

<sup>1</sup>The packet capture described in this analysis can be found at <https://github.com/brunchboy/dysentery/raw/master/doc/assets/powerup.pcapng>

<sup>2</sup>Values within packets are shown in hexadecimal, while packet lengths and byte offsets are discussed in decimal.



Figure 2: First-stage Mixer device number assignment packets

The tenth byte (inside what is labeled the header) is bolded because its value changes in the different types of packets which follow.

After about three of these packets are sent, another series of three begins. It is not clear what purpose these packets serve, because they are not yet asserting ownership of any device number; perhaps they are used when CDJs are powering up as part of the mechanism the mixer can use to tell them which device number to use based on which network port they are connected to?

In any case, these three packets have a data length of 44 bytes, are again sent to UDP port 50000 on the local network broadcast address, at roughly 300 millisecond intervals, and have the content shown in Figure 2.

The value  $N$  at byte 36 is 1, 2, or 3, depending on whether this is the first, second, or third time the packet is sent.

After these comes another series of three numbered packets. These appear to be claiming the device number for a particular device, as well as announcing the IP address at which it can be found. They have a data length of 50 bytes, and are again sent to UDP port 50000 on the local network broadcast address, at roughly 300 millisecond intervals, with the content shown in Figure 3.




Figure 3: Second-stage Mixer device number assignment packets

I identify these as claiming/identifying the device number because the value  $D$  at byte 46 is the same as the device number that the mixer uses to identify itself (0x21) and the same is true for the corresponding



Figure 4: Final-stage Mixer device number assignment packets

packets seen from the CDJs (they use device numbers 2 and 3, as they are connected to those ports/channels on the mixer).

As with the previous series of three packets, the value  $N$  at byte 47 takes on the values 1, 2, and 3 in the three packets.

These are followed by another three packets, perhaps the last stage of claiming the device number, again at 300 millisecond intervals, to the same port 50000. These shorter packets have 38 bytes of data and the content shown in Figure 4.

As before the value  $D$  at byte 36 is the same as the device number that the mixer uses to identify itself (0x21) and  $N$  at byte 37 takes on the values 1, 2, and 3 in the three packets.

Once those are sent, the mixer seems to settle down and send what looks like a keep-alive packet to retain presence on the network and ownership of its device number, at a less frequent interval. These packets are 54 bytes long, again sent to port 50000 on the local network broadcast address, roughly every second and a half. They have the content shown in Figure 5.




Figure 5: Mixer keep-alive packets

## 2 CDJ Startup

When a CDJ starts up the procedure and packets are nearly identical, with groups of three packets sent at 300 millisecond intervals to port 50000 of the local network broadcast address. The only difference

between Figure 6 and Figure 1 is the final byte, which is 0x01 for the CDJ, and was 0x02 for the mixer.




Figure 6: Initial announcement packets from CDJ

Similarly, the next series of three packets from the CDJ are nearly identical to those from the mixer. The only difference between Figure 7 and Figure 2 is byte 37 (immediately after the packet counter  $N$ ), which again is 0x01 for the CDJ, and was 0x02 for the mixer.

However it appears that in this capture the CDJ skips the second stage of claiming a device number, probably because it is configured to be automatically assigned a device number based on the port of the mixer to which it is connected, and we cannot see a packet that the mixer sent it assigning it that device number. Instead, it jumps right to the end of the third and final stage, sending a single 38-byte packet with header byte 10 set to tt 04 (which identified the three packets of the third stage when the mixer was starting up), with content identical to Figure 4.

Even though the value of  $N$  is 01, this is the only packet in this series that the CDJ sends. It would probably behave differently if configured to assign its own device number (behaving like we saw the mixer behave in claiming its device number).

The CDJ then moves to the keep-alive stage, sending out 54-byte packets with the content shown in Figure 8.

As seems to always be the case when comparing mixer and CDJ packets, the difference between this and Figure 5 is that byte 37 (following the device number  $D$ ) has the value 01 rather than 02, and the same is true of the second-to-last byte in each of the packets. (Byte 52




Figure 7: First-stage CDJ device number assignment packets




Figure 8: CDJ keep-alive packets

is 01 in Figure 8 and 02 in Figure 5.

### 3 Tracking BPM and Beats

For some time now, Afterglow<sup>3</sup> has been able to synchronize its light shows with music being played on Pioneer equipment by observing packets broadcast by the mixer to port 50001. Until recently, however, it was not possible to tell which player was the master, so there was no way to determine the down beat (the start of each measure). Now that it is possible to determine which CDJ is the master player using the packets described in Section 4, these beat packets have become far more useful, and Afterglow will soon be using them to track the down beat based on the beat number reported by the master player.

To track beats, open a socket and bind it to port 50001. The devices seem to broadcast two different kinds of packets to this port, a shorter packet containing 45 bytes of data, and a longer packet containing 96 bytes. The shorter packets seem to all have identical content, and do not seem to convey useful information, so we currently simply ignore them.

The 96-byte packets are sent on each beat, so even the arrival of the packet is interesting information, it means that the player is starting a new beat. (CDJs send these packets only when they are playing. The mixer sends them all the time, acting as a backup metronome when no other device is counting beats.) The content of these packets is shown in Figure 9.

The Device Number in *D* (bytes 33 and 95) is the Player Number as displayed on the CDJ itself, or 33 for the mixer, or another value for a computer running rekordbox.

The player's current pitch adjustment<sup>4</sup> can be found in bytes 85–87, labeled *Pitch*. It represents a three-byte pitch adjustment percentage,

<sup>3</sup><https://github.com/brunchboy/afterglow#afterglow>

<sup>4</sup>The mixer always reports a pitch of +0%.

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	51	73	70	74	31	57	6d	4a	4f	4c	<b>28</b>					
16	Device Name (padded with 00)															01
32	00	<i>D</i>	00	3c	00	00	01	00	00	00	00	00	00	00	00	00
48	00	00	00	00	00	00	00	00	00	00	00	ff	ff	ff	ff	ff
64	ff	ff	ff	ff	ff	ff	ff	ff	ff	ff	ff	ff	ff	ff	ff	ff
80	ff	ff	ff	ff	00	<i>Pitch</i>		00	00	<i>BPM</i>		<i>B<sub>b</sub></i>	00	00	<i>D</i>	

Figure 9: Beat packets

where 0x100000 represents no adjustment (0%), 0x000000 represents slowing all the way to a complete stop (−100%, reachable only in Wide tempo mode), and 0x200000 represents playing at double speed (+100%).

The pitch adjustment percentage represented by *Pitch* is calculated as follows:

$$100 \times \frac{(byte[85] \times 65536 + byte[86] \times 256 + byte[87]) - 1048576}{1048576}$$

The current BPM of the track playing on the device<sup>5</sup> can be found at bytes 90–91 (labeled *BPM*). It is a two-byte integer representing one hundred times the current track BPM. So, the current track BPM value to two decimal places can be calculated as:

$$\frac{byte[90] \times 256 + byte[91]}{100}$$

In order to obtain the actual playing BPM (the value shown in the BPM display), this value must be multiplied by the current pitch adjustment, calculated from *Pitch* as described above.<sup>6</sup>

The counter *B<sub>b</sub>* at byte 92 counts out the beat within each bar, cycling 1 → 2 → 3 → 4 repeatedly, and can be used to identify the down beat if it is coming from the master player.

## 4 Creating a Virtual CDJ

Although some useful information can be obtained simply by watching broadcast traffic on a network containing Pioneer gear, in order to

<sup>5</sup>The mixer passes along the BPM of the master player.

<sup>6</sup>Since the mixer always reports a pitch adjustment of +0%, its *BPM* value can be used directly without this additional step.

get important details it is necessary to cause the gear to send you information directly. This can be done by simulating a “Virtual CDJ”.<sup>7</sup>

To do this, bind a UDP server socket to port 50002 on the network interface on which you are receiving DJ-Link traffic, and start sending keep-alive packets to port 50000 on the broadcast address as if you were a CDJ. Follow the structure shown in Figure 8, but use the actual MAC and IP addresses of the network interface on which you are receiving DJ-Link traffic, so the devices can see how to reach you.

You can use a value like 5 for  $D$  (the device/player number) so as not to conflict with any actual players you have on the network, and any name you would like. As long as you are sending these packets roughly every 1.5 seconds, the other players and mixers will begin sending packets directly to the socket you have opened on port 50002.

Each device seems to send status packets roughly every 200 milliseconds.

We are just beginning to analyze all the information which can be gleaned from these packets, but here is what we know so far.<sup>8</sup>

#### 4.1 Mixer Status Packets

Packets from the mixer will have a length of 56 bytes and the content shown in Figure 10.

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	51	73	70	74	31	57	6d	4a	4f	4c	29					
16	Device Name (padded with 00)															01
32	00	<i>D</i>	00	14	<i>D</i>	00	00	d0	00	10	00	00	80	00	<i>BPM</i>	
48	00	10	00	00	00	09	00	<i>X</i>	<i>B<sub>b</sub></i>							

Figure 10: Mixer status packets

Packets coming from a DJM-2000 nexus connected as the only mixer on the network contain a value of 33 (0x21) for their Device Number  $D$  (bytes 33 and 36).

The current tempo in beats-per-minute identified by the mixer can be obtained as:

<sup>7</sup>Thanks are due to Diogo Santos for discovering the trick of creating a virtual CDJ in order to receive detailed status information from other devices.

<sup>8</sup>Examples of packets discussed in this section can be found in the capture at <https://github.com/brunchboy/dysentery/raw/master/doc/assets/to-virtual.pcapng>

$$\frac{\text{byte}[46] \times 256 + \text{byte}[47]}{100}$$

This value is labeled *BPM* in Figure 10. Unfortunately, this BPM seems to only be valid when a rekordbox-analyzed source is playing; when the mixer is doing its own beat detection from unanalyzed audio sources, even though it displays the detected BPM on the mixer itself, and uses that to drive its beat effects, it does not send that value in these packets.

The current beat number within a bar (1, 2, 3 or 4) is sent in *byte*[55], labeled *B<sub>b</sub>*. However, the beat number is *not* synchronized with the master player, and these packets do not arrive at the same time as the beat started anyway, so this value is not useful for much. The beat number should be determined, when needed, from beat packets (described in Section 3) that are sent by the master player.

The value at *byte*[54], labeled *X*, has an unknown meaning. It seems to start out with the value 0x00, and then change when a player starts playing to the value 0xff, but it may well do other things as well.

## 4.2 CDJ Status Packets

Packets from a CDJ will have a length of 212 bytes and the content shown in Figure 11.

The Device Number in *D* (bytes 33 and 36) is the Player Number as displayed on the CDJ itself. In the case of this capture, the CDJs were assigned Player Numbers 2 and 3.

The activity flag *A* at byte 39 seems to be 0 when the player is idle, and 1 when it is playing, searching, or loading a track.

The track number being played (its position within a playlist or other scrolling list of tracks, as displayed on the CDJ) can be found at bytes 50 and 51, labeled *Track*. (It may be a 4-byte value and also include bytes 48 and 49, but that would seem an unmanageable number of tracks to search through.)

There are a number of bytes, labeled *t<sub>2</sub>* through *t<sub>7</sub>*, whose purpose is as yet undetermined. They are all zero when there is no track loaded, but take different values when a track is loaded. (Actually, there are many other bytes than these which behave like this, so they will probably be removed from the chart in a future version of this document.)

Byte 106, labeled *U<sub>a</sub>* (for “USB activity”), alternates between the values 4 and 6 when there is USB activity—it may even alternate in time with the flashing USB indicator LED on the player. Byte 111 (*U<sub>l</sub>* for “USB local”) has the value 4 when there is no USB media loaded,



	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	51	73	70	74	31	57	6d	4a	4f	4c	0a					
16	Device Name (padded with 00)															
32	03	<i>D</i>	00	b0	<i>D</i>	00	01	<i>A</i>	00	00	00	00	00	00	00	00
48	00	00	<i>Track</i>		00	00	<i>t</i> <sub>2</sub>	<i>t</i> <sub>3</sub>	00	00	<i>t</i> <sub>4</sub>	<i>t</i> <sub>5</sub>	00	00	00	00
64	00	00	00	00	00	00	<i>t</i> <sub>6</sub>	<i>t</i> <sub>7</sub>	00	00	00	00	00	00	00	00
80	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
96	00	00	00	00	00	00	00	00	01	00	<i>U</i> <sub>a</sub>	04	00	00	00	<i>U</i> <sub>l</sub>
112	00	00	00	04	00	<i>U</i> <sub>g</sub>		00	00	01	00	00	<i>P</i> <sub>1</sub>	31	2e	32 34
128	00	00	00	00	00	00	<i>Sync</i> <sub><i>n</i></sub>		00	<i>F</i>	ff	<i>P</i> <sub>2</sub>	00	<i>Pitch</i> <sub>1</sub>		
144	<i>l</i> <sub>1</sub>		<i>BPM</i>		7f	ff	ff	ff	00	<i>Pitch</i> <sub>2</sub>			00	<i>P</i> <sub>3</sub>	00	ff
160	<i>Beat</i>				<i>Cue</i>		<i>B</i> <sub><i>b</i></sub>	00	00	00	00	00	00	00	00	00
176	00	00	00	00	00	00	10	00	00	00	00	00	00	00	00	00
192	00	<i>Pitch</i> <sub>3</sub>			00	<i>Pitch</i> <sub>4</sub>			<i>Packet</i>				0f	00	00	00
208	00	00	00	00												

Figure 11: CDJ status packets

7	6	5	4	3	2	1	0
1	Play	Master	Sync	On-Air	1	0	0

Figure 12: CDJ state flag bits

0 when USB is loaded, and 2 when the USB Stop button has been pressed and the USB media is being unmounted.

Byte 117, labeled  $U_g$  (for “USB global”), appears to have the value 1 whenever USB media is present in any player on the network, whether or not the Link option is chosen in the other players, and 0 otherwise. I don’t know if that is true of other linkable SD media as well; this needs more investigation.

Byte 123, labeled  $P_1$ , appears to describe the current play mode. The values that have been seen so far, and their apparent meanings, are:

- 0 No track is loaded.
- 3 The player is playing normally.
- 4 The player is playing a loop.
- 5 The player is paused anywhere other than the cue point.
- 6 The player is paused at the cue point.
- 9 The player is searching forwards or backwards.
- 17 The player reached the end of the track and stopped.

The value  $Sync_n$  at bytes 134–135 seems to increment whenever the player syncs to a new tempo master (another player or the mixer). I am assuming it is just a 2-byte value, because I tried syncing 256 times and saw the counter expand from byte 135 to include byte 134. It may actually be a 4-byte value and also involve bytes 132 and 133, but I wasn’t going to try changing sync 65,536 times or more to find out. Perhaps we could write software to test that someday by forcing tempo master changes.

Byte 137, labeled  $F$ , is a bit field containing some very useful state flags, detailed in Figure 12.<sup>9</sup>

Byte 139, labeled  $P_2$  seems to be another play state indicator, having the value 122 when playing and 126 when stopped. When the CDJ is trying to play, but is being held in place by the DJ holding down on the jog wheel,  $P_1$  considers it to be playing (value 3), while  $P_2$  considers it to be stopped (value 126).

<sup>9</sup>We have not yet seen any other values for bits 0, 1, 2, or 7 in  $F$ , so we’re unsure if they also carry meaning. If you ever find different values for them, please let us know by filing an Issue at <https://github.com/brunchboy/dysentery/issues>

There are four different places where pitch information appears in these packets: *Pitch*<sub>1</sub> at bytes 141–143, *Pitch*<sub>2</sub> at bytes 153–155, *Pitch*<sub>3</sub> at bytes 193–195, and *Pitch*<sub>4</sub> at bytes 197–199.

Each of these values represents a three-byte pitch adjustment percentage, where 0x100000 represents no adjustment (0%), 0x000000 represents slowing all the way to a complete stop (−100%, reachable only in Wide tempo mode), and 0x200000 represents playing at double speed (+100%).

Note that if playback is stopped by pushing the pitch fader all the way to −100% in Wide mode, both *P*<sub>1</sub> and *P*<sub>2</sub> still show it as playing, which is different than when the jog wheel is held down, since *P*<sub>2</sub> shows a stop in the latter situation.

Here is how the pitch adjustment percentage represented by *Pitch*<sub>1</sub> would be calculated:

$$100 \times \frac{(\text{byte}[141] \times 65536 + \text{byte}[142] \times 256 + \text{byte}[143]) - 1048576}{1048576}$$

We don't know why there are so many copies of the pitch information, or all circumstances under which they might differ from each other, but it seems that *Pitch*<sub>1</sub> and *Pitch*<sub>3</sub> report the current pitch adjustment actually in effect (as reflected on the BPM display), whether it is due to the local pitch fader, or a synced tempo master.

*Pitch*<sub>2</sub> and *Pitch*<sub>4</sub> are always tied to the position of the local pitch fader, unless Tempo Reset is active, effectively locking the pitch fader to 0% and *Pitch*<sub>2</sub> and *Pitch*<sub>4</sub> to 0x100000, or the player is paused or the jog wheel is being held down, freezing playback and locking the local pitch to −100%, in which case they both have the value 0x000000.

When playback stops, either due to the play button being pressed or the jog wheel held down, the value of *Pitch*<sub>4</sub> drops to 0x000000 instantly, while the value of *Pitch*<sub>2</sub> drops over time, reflecting the gradual slowdown of playback which is controlled by the player's brake speed setting. When playback starts, again either due to the play button being pressed or the jog wheel being released, both *Pitch*<sub>2</sub> and *Pitch*<sub>4</sub> gradually rise to the target pitch, at a speed controlled by the player's release speed setting.

If the player is *not* synced, but the current pitch is different than what the pitch fader would indicate (in other words, the player is in the mode where it tells you to move the pitch fader to the current BPM in order to change the pitch), moving the pitch fader changes the values of *Pitch*<sub>2</sub> and *Pitch*<sub>4</sub> until they match *Pitch*<sub>1</sub> and *Pitch*<sub>3</sub> and begin to affect the actual effective pitch. From that point on, moving the pitch fader sets the value of all of *Pitch*<sub>1</sub>, *Pitch*<sub>2</sub>, *Pitch*<sub>3</sub>, and *Pitch*<sub>4</sub>. This all seems more complicated than it really needs to be...

The current BPM of the track (the BPM at the point that is currently being played, or at the location where the player is currently paused) can be found at bytes 146–147 (labeled *BPM*). It is a two-byte integer representing one hundred times the current track BPM. So, the current track BPM value to two decimal places can be calculated as:

$$\frac{\text{byte}[146] \times 256 + \text{byte}[147]}{100}$$

In order to obtain the actual playing BPM (the value shown in the BPM display), this value must be multiplied by the current effective pitch, calculated from  $Pitch_1$  as described above.

Because Rekordbox and the CDJs support tracks with variable BPM, this value can and does change over the course of playing such tracks. When no track is loaded, *BPM* has the value 0xffff.

The meaning of value  $l_1$  (bytes 144–145) is not currently known. It may simply reflect whether a track is loaded or not: it seems to have the value 0x7fff when no track is loaded, and the value 0x8000 when a track is loaded.

Byte 157 (labeled  $P_3$ ) seems to communicate additional information about the current play mode, with the following meanings that we have found so far:

- 0 No track is loaded.
- 1 The player is paused or playing in Reverse mode.
- 9 The player is playing in Forward mode with jog mode set to Vinyl.
- 13 The player is playing in Forward mode with jog mode set to CDJ.

The 4-byte beat counter (which counts each beat from 1 through the end of the track) is found in bytes 160–163, labeled *Beat*. When the player is paused at the start of the track, this seems to hold the value 0, even though it is beat 1, and when no track is loaded, this holds the value 0xffffffff.

The counter  $B_b$  at byte 166 counts out the beat within each bar, cycling  $1 \rightarrow 2 \rightarrow 3 \rightarrow 4$  repeatedly, and can be used to identify the down beat (as is used in the Master Player display on the CDJs as a mixing aid). Again, when no track is loaded, this holds the value 0. If you want to synchronize events to the down beat, use the CDJ status packets'  $F$  value to identify the master player, but use the beat packets sent by that player (described in Section 3) to determine when the beats are actually happening.

A countdown timer to the next saved cue point is available in bytes 164–165 (labeled *Cue*). If there is no saved cue point after the current play location in the track, or if it is further than 64 bars ahead, these bytes contain the value 0x01ff and the CDJ displays “-.- bars”.

As soon as there are just 64 bars (256 beats) to go before the next cue point, this value becomes 0x0100. This is the point at which the CDJ starts to display a countdown, which it displays as “63.4 bars”. As each beat goes by, this value decreases by 1, until the cue point is about to be reached, at which point the value is 0x0001 and the CDJ displays “0.1 bars”. On the beat on which the cue point was saved the value is 0x0000 and the CDJ displays “0.0 Bars”. On the next beat, the value becomes determined by the next cue point (if any) in the track.

Bytes 200-203 seem to contain a 4-byte packet counter labeled *Packet*, which is incremented for each packet sent by the player. (I am just guessing it is four bytes long, I have not yet watched long enough for the count to need more than the last three bytes).

This analysis is incomplete; there are many bytes whose purpose is not yet known, not all of which are even mentioned here yet.

## 5 What's Missing?

We know this analysis isn't complete. There are some values we know are available and useful (because CDJs and rekordbox itself show them), but we have not figured out how to find them yet.

### 5.1 Track Metadata

The biggest piece of information that would be extremely helpful for a light show operator to be able to see, and which we have not yet figured out how to obtain, is the name and artist of the track playing on a player.

While it may be that the rekordbox track ID is present in the status packets, we very much hope to avoid having to mount the rekordbox database over NFS and figure out how to read it. That is probably a few steps too far in complexity and risk.

In the hope we can avoid that, we've ordered an Ethernet switch that offers port mirroring, so we can get a look at all the traffic between two physical CDJs to see what happens when the Link Info page is pulled up on one of them.

### 5.2 CDJ Packets to Rekordbox

Performing a packet capture while rekordbox is running reveals that the CDJs send unicast packets to rekordbox address on port 50000, in addition to the packets they normally broadcast on that port. Figuring out how to pose as rekordbox might be useful in order to see what

additional data these can offer, although that may be much more work than posing as a CDJ.

### 5.3 Dysentery

If you have access to Pioneer equipment and are willing to help us validate this analysis, and perhaps even figure out more details, you can find the tool that is being used to perform this research at:

<https://github.com/brunchboy/dysentery>



<http://deepsymmetry.org>